

ENV 710: Lecture 19

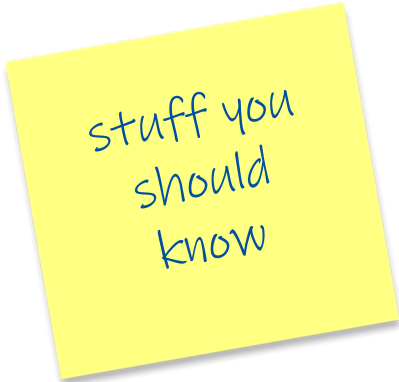
binomial logistic regression

generalized linear models

binomial logistic regression

learning goals

- binomial logistic regression
 - grouped responses and modeling proportions
 - odds, log-odds, probabilities
 - what is it? when to use it?
 - how does it work?
 - interpretation of coefficients



stuff you
should
know

logistic regression

- binary response variable
 - success or failure (1 or 0)
 - model the probability of success (π)

$$\pi = \beta_0 + \beta_1 X_1$$

- model log odds

$$\text{logit}(\pi) = \log[\pi/(1 - \pi)]$$


$$\text{logit}(\pi) = \beta_0 + \beta_1 X_k + \dots + \beta_k X_k$$

- calculate probabilities


$$\pi = \frac{\exp(\beta_0 + \beta_1 X_k + \dots + \beta_k X_k)}{1 + \exp(\beta_0 + \beta_1 X_k + \dots + \beta_k X_k)}$$

goal is to model the probability of 'success'

transform response to be linearly related to the IV's through the logit transform



back transform to probability using the inverse logit



logistic regression for binomial counts

- **binomial count** = sum of independent binary responses
- $Y \sim \text{binomial}(n, \pi)$, with population proportion, π , and n trials
- observed proportion of 1's is Y/n – binomial proportion
- continuous proportions of amounts cannot be modeled
 - e.g., the proportion of fat that is saturated fat
 - numerator and denominator are not integers, and no n is involved
 - when response variable is a continuous proportion, use ordinary regression methods

- model how population proportion depends on the explanatory variables through a nonlinear link function

$$\text{logit}(\pi) = \beta_0 + \beta_1 x_1 + \beta_2 x_2$$

- interpretation of model is same as for binary logistic regression, which is a special case of the binomial model in which all n_i 's are 1

example

beetles

Beetles are exposed to different doses of carbon disulphide (mf/L) for 5 hours.
Does dosage level affect mortality?

dose-response curve

| Dose | # Exposed | # Killed | Proportion |
|------|-----------|----------|------------|
| 49.1 | 59 | 6 | 0.102 |
| 53.0 | 60 | 13 | 0.217 |
| 56.9 | 62 | 18 | 0.290 |
| 60.8 | 56 | 28 | 0.500 |
| 64.8 | 63 | 52 | 0.825 |
| 68.7 | 59 | 53 | 0.898 |
| 72.6 | 62 | 61 | 0.985 |
| 76.5 | 60 | 60 | 1.000 |

- lowest dose, $n = 59$ beetles, $k = 6$
- for each beetle, outcome is death or survival
 - death = success
 - p = probability that an individual beetle will die after the dose
 - probability that y assumes a value k is given by:

$$P(y = k) = \binom{n}{k} p^k (1 - p)^{n-k}$$

example

beetles

```
b1 <- glm(cbind(killed, exposed - killed) ~ dose,  
          family = binomial)  
summary(b1)
```

```
glm(formula = cbind(killed, alive) ~ dose,  
     family = binomial)
```

Coefficients:

| | Estimate | Std. Error | z value | Pr(> z) |
|-------------|-----------|------------|---------|------------|
| (Intercept) | -14.82300 | 1.28959 | -11.49 | <2e-16 *** |
| dose | 0.24942 | 0.02139 | 11.66 | <2e-16 *** |

(Dispersion parameter for binomial family taken to be 1)


Null deviance: 284.2024 on 7 degrees of freedom
Residual deviance: 7.3849 on 6 degrees of freedom
AIC: 37.583

```
betas <- coef(b1)
```

```
exp(betas)  
(Intercept)      dose  
3.651351e-07 1.283275e+00
```

```
exp(cbind(OR = coef(b1), confint(b1)))  
  
              OR          2.5 %          97.5 %  
(Intercept) 3.651351e-07 2.530115e-08 4.02275e-06  
dose        1.283275e+00 1.233305e+00 1.34146e+00
```

dose effect plot




example

beetles



```
b1 <- glm(cbind(killed, exposed - killed) ~ dose,  
          family = binomial)  
  
plot(dose, killed/exposed, xlab="Dose, mf/L",  
      ylab="Probability of beetle mortality",  
      ylim=c(0,1),  
      pch = 21, bg = "black")  
  
x <- seq(min(dose), max(dose), length = 100)  
  
lines(x, predict(b1, data.frame(dose=x),  
                      type="response"), lwd = 2, type = "l",  
      col = "red")  
  
# predict specific values  
predict(b1, data.frame(dose = c(50, 65, 75)), type =  
"response")
```

```
1          2          3  
0.08688964 0.80043514 0.97982979
```






example

beetles



```
plot(cooks.distance(b1))
```



goodness-of-fit

deviance test for lack-of-fit

- significant p-value suggests that a residual deviance as large or larger than that observed under the model is highly unlikely
 - model is not adequate or lack of fit
- deviance**: $-2 \log(\text{likelihood})$ evaluated at the maximum likelihood of the parameters
- change in deviance has a Chi-squared distribution with df's equal to the change in number of parameters in the model

```
1-pchisq(b1$dev, b1$df.resid)
[1] 0.2867122
```

```
pchisq(b1$dev, b1$df.resid, lower.tail = F)
[1] 0.2867122
```

pseudo- R^2

```
require(psc1)
pR2(b1)
      llh      llhNull      G2
-16.7914741 -155.2002438  276.8175395

McFadden      r2ML      r2CU
0.8918077      1.0000000      1.0000000
```

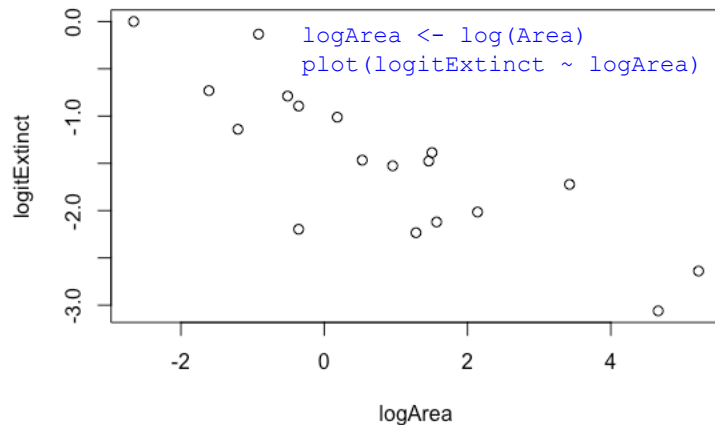
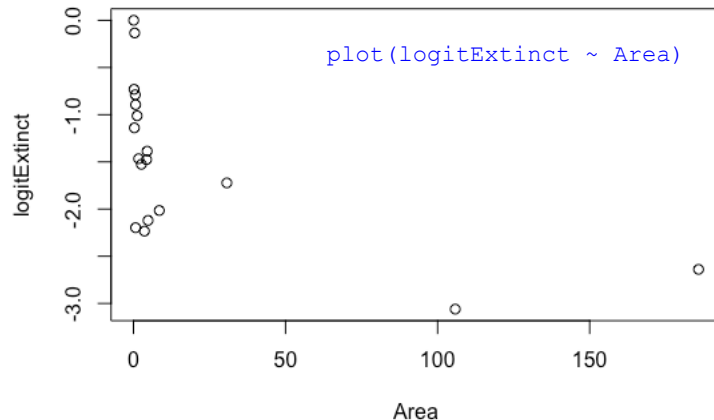
example

Kruunitt Island Archipelago

| | Island | Area | AtRisk | Extinct |
|---|----------------|-------|--------|---------|
| 1 | Ulkokruunni | 185.8 | 75 | 5 |
| 2 | Maakruunni | 105.8 | 67 | 3 |
| 3 | Ristikari | 30.7 | 66 | 10 |
| 4 | Isonkivenletto | 8.5 | 51 | 6 |
| 5 | Hietakraasukka | 4.8 | 28 | 3 |
| 6 | Kraasukka | 4.5 | 20 | 4 |

- logit of the probability of extinctions is assumed to be linear in $X = \log(\text{Area})$
- logistic regression on binary responses is a special case of the binomial model in which all n_i 's are 1

```
proportionExtinct <- Extinct/AtRisk  
oddsExtinct <- proportionExtinct/(1 - proportionExtinct)  
logitExtinct <- log(oddsExtinct) # Logit = Log Odds
```



example

Krunnit Island Archipelago

```
k2 <- glm(cbind(Extinct, AtRisk - Extinct) ~  
          log(Area), family = binomial)  
summary(k2)
```

Coefficients:

| | Estimate | Std. Error | z value | Pr(> z) |
|-------------|----------|------------|---------|--------------|
| (Intercept) | -1.19620 | 0.11845 | -10.099 | < 2e-16 *** |
| log(Area) | -0.29710 | 0.05485 | -5.416 | 6.08e-08 *** |


(Dispersion parameter for binomial family taken to be 1)

Null deviance: 45.338 on 17 degrees of freedom
Residual deviance: 12.062 on 16 degrees of freedom
AIC: 75.394

```
betas <- coef(k2)  
odds <- exp(betas)  
odds
```

| (Intercept) | log(Area) |
|-------------|-----------|
| 0.3023423 | 0.7429670 |

for every 1 unit increase in
log(Area), odds of extinction
is 74.2% that of the previous
area, or a reduction of 25.8%
in odds of extinction



goodness of fit

```
pchisq(k2$deviance, k2$df.residual, lower.tail = F)
[1] 0.7397351
```

```
pR2(k2) [4]
```

```
fitting null model for pseudo-r2
McFadden
0.3179176
```





```
knnull <- update(k2, .~.-log(Area))
lrtest(knull, k2)
```

Likelihood ratio test

```
Model 1: cbind(Extinct, AtRisk - Extinct) ~ 1
Model 2: cbind(Extinct, AtRisk - Extinct) ~ log(Area)
```

| | #Df | LogLik | Df | Chisq | Pr(>Chisq) |
|---|-----|---------|----|--------|---------------|
| 1 | 1 | -52.335 | | | |
| 2 | 2 | -35.697 | 1 | 33.276 | 7.994e-09 *** |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1



example

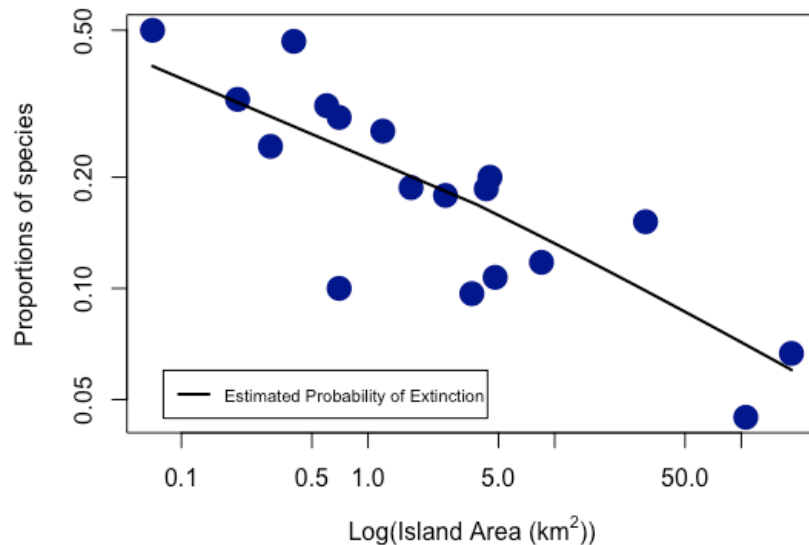
Krunnit Island,

```
1-exp(betas[2]*log(2))  
log(Area)  
0.1861153
```

```
1-2^confint(k2, 2)  
      2.5 %      97.5 %  
0.2462041 0.1247743
```

$$\ln(x^y) = y \cdot \ln(x)$$

- with each doubling of island area, there is a ~ 19% reduction in the odds of extinction
- each halving of area is associated with a 23% increase in extinction



```
plot(proportionExtinct ~ Area, log="xy",  
      ylab="Proportions of species",  
      xlab=expression(paste("Log(Island Area (km^2)", "))),  
      pch=21, lwd=2, bg="darkblue", col="darkblue", cex=2)  
  
dummyArea <- seq(min(Area), max(Area), length=50)  
lp <- betas[1] + betas[2]*log(dummyArea)  
myProbability <- exp(lp)/(1 + exp(lp))  
  
lines(myProbability ~ dummyArea, lwd=2)
```



Questions?